

AI Data Centers and Power Grids: Understanding and Possible Problem

Richard Harnisch

July 23, 2026

1. Overall Understanding

AI training clusters invalidate the load-diversity assumption because they behave as one synchronized, geographically concentrated load. The papers explore different layers of this problem, from workload and infrastructure simulation to facility design, operational flexibility, fault response, and market coordination. A central gap is therefore the absence of validated interfaces that connect decisions at the compute layer to facility-level power behavior and grid-level consequences.

2. Personal Thoughts

I am most drawn to the approach taken by Phythesis. My work with LLM harnesses and agentic validation has been very engaging in the past. I also find the approach of using gradient descent on the component parameters followed by proximity matching to the closest real-world component to be novel. I would be interested in expanding this, potentially by integrating the LLM-created design process with one of the programmatic simulators in the other papers.

I found the demand response paper both interesting and engaging. I think the market-based approach using price signals is smart because it can operate within the existing regulatory framework instead. This is preferable because legislation is slow and subject to political lobbying. Instead, market actors can be directly compensated for grid-supporting behavior, incentivizing stability-supporting behavior.

I found that “From Barrier to Bridge” makes a valid argument for data-center-grid codesign, but it was sometimes difficult to follow their points and claims, for example about regulations. In essence, their main point (as I understand it) is that power grid operators and data center scalers should work more closely to avoid grid instability. However, the real question and contribution (to me) would be how we can incentivize this while minimizing costs to the end user through higher power prices.

ASTRA-SIM 3.0 was technically complex. I would need more time to read into the prerequisite knowledge, as I have no experience with GPU components. However, they make an interesting contribution. The InfraGraph framework seems most directly linked to the thread formed by the other papers, so further exploration on my part would likely focus on this interconnection.

DCGen provides a fast way to obtain a datacenter configuration based on a set of requirements, including into the future. However, it seems to be limited in terms of making changes to the generated design. To me it seems the most straightforward use seems to be its ability to forecast how different constraints on data centers will behave into the future (2027, 2029). However, it may be an interesting direction to use the DCGen-generated layouts as a starting point for LLM or other forms of refinement and further adaptation. Another potential direction could be allowing the DCGen-created designs to be tested by simulations. I anticipate this to be mostly a compatibility-engineering problem.

“Low-voltage ride-through” proposes an engineering approach to the grid-stability problem where the demand reduction paper proposes a market-based approach. The two complement each other in that the price mechanisms described in the latter can incentivize adoption of solutions described in the former. I see these two papers as complementary, with the caveat that low-voltage ride-through assumes requiring compliance through regulatory means rather than using price signals. However, a price signal could be used to incentivize committing to specific ride-through requirements.

3. Possible Research Avenues

- **LLM-simulator integration:** Extend Phythesis by connecting its LLM-generated designs and physics-guided refinement to the programmatic simulators from the other papers. Potentially use DCGen-generated configurations as a starting point for LLM-based refinement and adaptation to new requirements, or modelling how well the proposed design can comply with voltage-ride-through designs. The contribution might be a similar system that generates layouts which can comply with any one of a set of ride-through requirement specifications. I would favor one of these directions.
- **Extending InfraGraph to interface with Phythesis:** Similarly, explore whether ASTRA-SIM’s InfraGraph framework is integrable for Phythesis. This would mean that phythesis would additionally deliver an InfraGraph representation of the data center, allowing it to interface directly with ASTRA-SIM.
- **Expanding DCGen design evaluation:** Test DCGen-generated designs with the simulation and control approaches from the other papers, focusing on the compatibility engineering needed between their representations. Potentially expand DCGen to consider or estimate data center cost under different grid stability and compensation conditions.
- **Market incentives for grid support:** Further develop price-based mechanisms that reward data centers for flexible, grid-supporting behavior while limiting cost increases for end users. Potentially explore how the adoption of the solutions proposed in the low-voltage ride-through paper can be further incentivized.